

CS289–Spring 2017 — Homework 6 Solutions

Shuhui Huang, SID 3032129712

1. Problem 1: Derivations

We have two weight matrices.

V is a $785 \times 200(n_{in} + 1) \times n_{hid}$ matrix that connects the input units to the hidden units. The last column of V is made up of bias terms, which are multiplied by 1 instead of an input value.

W is a $201 \times 26(n_{hid} + 1) \times n_{out}$ matrix that connects the hidden units to the output units. The last column is bias terms.

Let X be the input layer, h be the hidden layer, Z be the output layer, y is the ground truth label So we have :

$$h_j = \tanh(\sum_i V_{ij} X_i), Z_k = \text{sigmoid}(\sum_j W_{jk} h_j)$$

As we use the cross-entropy loss function:

$$L(z, y) = -\sum_{j=1}^{26} y_j \ln Z_j + (1 - y_j) \ln(1 - Z_j)$$

$$\text{So } \frac{\partial L}{\partial Z_k} = -\left(\frac{y_k}{Z_k} - \frac{1-y_k}{1-Z_k}\right) = \frac{Z_k - y_k}{Z_k(1-Z_k)}$$

$$\frac{\partial L}{\partial W_{jk}} = \frac{\partial L}{\partial Z_k} \frac{\partial Z_k}{\partial W_{jk}} = \frac{Z_k - y_k}{Z_k(1-Z_k)} Z_k(1 - Z_k) h_j = (Z_k - y_k) h_j$$

$$\frac{\partial L}{\partial V_{ij}} = \left(\sum_k \frac{\partial L}{\partial Z_k} \frac{\partial Z_k}{\partial h_j}\right) \frac{\partial h_j}{\partial V_{ij}} = \left(\sum_k (Z_k - y_k) W_{jk}\right) (1 - h_j^2) X_i$$

\therefore we have:

$$\frac{\partial L}{\partial W} = h(Z - y)^T$$

$$\frac{\partial L}{\partial V} = X[W(Z - y)(1 - h^2)]^T$$

Which is the partial derivatives of L with respect to V and W we want to get.

2.Problem 2: Implementation

(a) Any hyperparameters that you tuned.

I first used learning rate=0.001, then tuned it to 0.05,max iteration=500000, loss function is cross entropy error,and hidden layer have 200 units.


(b) my training accuracy.

0.8823

(c) my validation accuracy:

0.8656

(d) A plot of the loss value versus the number of iterations. You may sample (i.e., compute the loss every 1000 iterations).





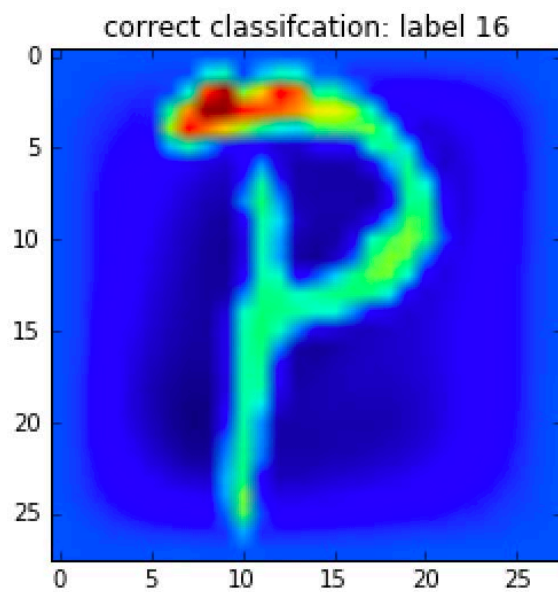


(e) my Kaggle score: 0.86115

display name:Shuhui Huang

(f) My code is attached as an appendix at the end.


3.Problem 3: Visualization

5 digits (and their labels) from the validation set that your neural network correctly classifies:

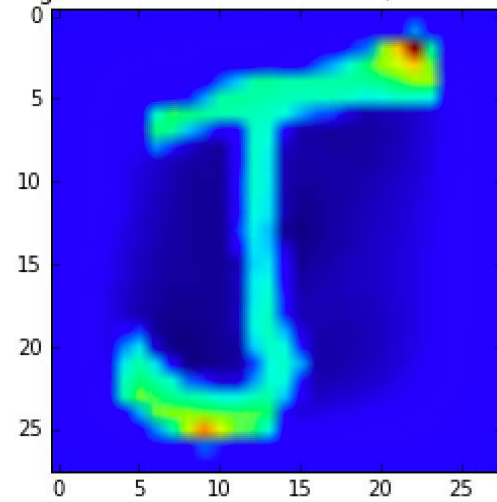


5 digits (and their labels) from the validation set that my neural network does not correctly classify:


wrong classification: true label 2, classified as 5




wrong classification: true label 10, classified as 9




wrong classification: true label 15, classified as 14



wrong classification: true label 9, classified as 12



wrong classification: true label 17, classified as 15



4.Problem 4: Bells and Whistles

I changed the number of hidden layer units from 200 to 500, and also changed the initialization of weights from a Gaussian distribution with mean 0 and variance $\sigma^2 \propto 1/\eta$, where η is the fan-in of the neuron the weight is an input to.

In addition, I also use different learning rates for different layers, and have those rates decay over time. learning rate=learning rate $\ast \frac{k}{k+1}$, where k is $\frac{\text{iteration numbers}}{\text{sample numbers}}$.

My code has been attached in the appendix.

```
#!/usr/bin/env python3
```

```
# -*- coding: utf-8 -*-
```

```
"""
```

```
Created on Tue Mar 28 10:20:16 2017
```

```
@author: huangshuhui
```

```
"""
```

```
#####problem2#####
```

```
import numpy as np
```

```
import scipy.io
```

```
import random
```

```
from sklearn.preprocessing import StandardScaler
```

```
from scipy.special import expit
```

```
import matplotlib.pyplot as plt
```

```
import math
```

```
#load the data#
```

```
data_raw=scipy.io.loadmat('/Users/huangshuhui/Google  
Drive/study/cs289/hw2017/hw6/hw6_data_dist/letters_data.mat')
```

```
test_x=data_raw['test_x']
```

```
train_x=data_raw['train_x']
```

```
train_label=data_raw['train_y']
```

```
#normalize and shuffle the data#
```

```
scaler = StandardScaler()
```

```
scaler.fit(train_x)
```

```
train_x = scaler.transform(train_x)
```

```
test_x = scaler.transform(test_x)
```

```
x_y = list(zip(train_x, train_label))
```

```
random.shuffle(x_y)
```

```
train_x = np.array([e[0] for e in x_y])
```

```
train_label = np.ravel([e[1] for e in x_y])
```

```
train_y = np.zeros(shape=(train_x.shape[0],27))
```

```
for i in range(len(train_y)):
```

```
    train_y[i][train_label[i]]=1
```

```
#the original neural network#
```

```
class NN:
```

```
    def __init__(self,learning_rate,max_iteration,cost_function="cross-  
entropy",regularizer=0,n_hid=200):
```

```
        self.learning_rate = learning_rate
```

```
        self.max_iteration = max_iteration
```

```
        self.cost_function = cost_function
```

```
self.regularizor = regularizor
self.n_hid = n_hid
```

```
def train(self,train_x,train_y):
```

```
    ##### initialize all weights randomly #####
```

```
    n_in = train_x.shape[1]
```

```
    n_hid = self.n_hid
```

```
    n_out = train_y.shape[1]
```

```
    self.V = 0.01*np.random.randn(n_in+1,n_hid)
```

```
    self.W = 0.01*np.random.randn(n_hid+1,n_out)
```

```
    self.iteration_nums = []
```

```
    self.error_ls = []
```

```
    self.accuracy_ls = []
```

```
    ##### iteration #####
```

```
    t=0
```

```
    while(t<self.max_iteration):
```

```
    # while(t<self.max_iteration):
```

```
        t+=1
```

```
        ### pick a random point and add bias term ###
```

```
        rand_index = np.random.randint(0,len(train_y))
```

```
        x = np.append(train_x[rand_index],1)
```

```
        y = train_y[rand_index]
```

```
        x = np.array([x]).T
```

```
        y = np.array([y]).T
```

```
        ### forward pass ###
```

```
        z = np.tanh((self.V.T).dot(x))
```

```
        z = np.vstack((z,[1]))
```

```
        h = expit((self.W.T).dot(z))
```

```
        ### backward pass ###
```

```
        if self.cost_function == "cross-entropy":
```

```
            delta_V,delta_W = self.Cross_entropy_delta_W(z,h,y,x,self.W)
```

```
        ### gradient descent update ###
```

```
        self.V = (1-self.regularizor)*self.V - self.learning_rate*delta_V
```

```
        self.W = (1-self.regularizor)*self.W - self.learning_rate*delta_W
```

```
    ### total training error and classification accuracy vs. iteration ###
```

```
    if t%1000==0:
```

```
        self.iteration_nums.append(t)
```

```
        if self.cost_function=="cross-entropy":
```

```
            error = self.cross_entropy_error(train_x,train_y,self.V,self.W)
```

```
            predict_labels = self.predict(train_x)
```

```

        true_labels = train_y.argmax(axis=1)
        accuracy = sum(predict_labels==true_labels)/len(true_labels)
        print(t, " : ",error,accuracy)
        self.error_ls.append(error)
        self.accuracy_ls.append(accuracy)

def Cross_entropy_delta_W(self,z,h,y,x,W):
    delta_W = z.dot((h-y).T)
    delta_V = x.dot(((W[0:-1].dot(h-y))*(1-(z[0:-1])**2)).T)
    return delta_V,delta_W

def cross_entropy_error(self,train_x,train_y,V,W):
    ### calculate total cross entropy error ###
    X = np.column_stack((train_x, np.ones(train_x.shape[0])))
    Z = np.tanh(X.dot(V))
    Z = np.column_stack((Z, np.ones(Z.shape[0])))
    H = 1/(np.exp(-Z.dot(W))+1)
    Cross_entropy_error = np.sum(-(train_y*np.log(H)+(1-train_y)*np.log(1-H)))
    return Cross_entropy_error

def predict(self,test_x):
    X = np.column_stack((test_x, np.ones(test_x.shape[0])))
    Z = np.tanh(X.dot(self.V))
    Z = np.column_stack((Z, np.ones(Z.shape[0])))
    H = Z.dot(self.W)
    predicted_labels = H.argmax(axis=1)
    return predicted_labels

#test the training and validation accuracy#
sub_train_x = train_x[0:100000]
sub_train_y = train_y[0:100000]
sub_validate_x = train_x[100000:120000]
sub_validate_y = train_y[100000:120000]
test_NN = NN(learning_rate=0.01,max_iteration=600000,cost_function="cross-entropy",regularizer=0,n_hid=500)
test_NN.train(sub_train_x,sub_train_y)
predict_y = test_NN.predict(sub_validate_x)
true_y = sub_validate_y.argmax(axis=1)
validate_accuracy = sum(predict_y==true_y)/len(true_y)
validate_accuracy

#plot iteration and loss#
cross_NN = NN(learning_rate=0.001,max_iteration=500000,cost_function="cross-entropy",regularizer=0,n_hid=200)

```



```
cross_NN.train(train_x,train_y)
plt.plot(cross_NN.iteration_nums, cross_NN.error_ls, 'r-')
plt.xlabel('iteration numbers')
plt.ylabel('loss value')
plt.title('loss value versus number of iterations')
```

```
#####problem3#####
```

```
##Visualization
```

```
correct=0
```

```
wrong=0
```

```
i=0
```

```
label1,label2,label3=[],[],[]
```

```
validate_correct=[]
```

```
validate_wrong=[]
```

```
while correct<5:
```

```
    if predict_y[i]==true_y[i] and true_y[i] not in label1:
```

```
        correct+=1
```

```
        label1.append(true_y[i])
```

```
        validate_correct.append(sub_validate_x[i])
```

```
    i+=1
```

```
i=0
```

```
while wrong<5:
```

```
    if predict_y[i]!=true_y[i] and true_y[i] not in label2:
```

```
        wrong+=1
```

```
        label2.append(true_y[i])
```

```
        label3.append(predict_y[i])
```

```
        validate_wrong.append(sub_validate_x[i])
```

```
    i+=1
```

```
a=validate_wrong[4]
```

```
plt.imshow(a.reshape(28, 28))
```

```
plt.title('wrong classfication: true label 17, classified as 15' )
```

#####problem4#####

#neural network(Bells and Whistles)#

class NN2:

def __init__(self,learning_rate,max_iteration,cost_function="cross-entropy",regularizor=0,n_hid=200):

self.learning_rate = learning_rate

self.max_iteration = max_iteration

self.cost_function = cost_function

self.regularizor = regularizor

self.n_hid = n_hid

def train(self,train_x,train_y):

initialize all weights randomly

n_in = train_x.shape[1]

n_hid = self.n_hid

n_out = train_y.shape[1]

self.V =

np.vstack((math.sqrt(2.0/(n_in+n_hid))*np.random.randn(n_in,n_hid),0.01*np.ones(n_hid)))

self.W =

np.vstack((math.sqrt(2.0/(n_hid+n_out))*np.random.randn(n_hid,n_out),0.01*np.ones(n_out))
)

self.iteration_nums = []

self.error_ls = []

self.accuracy_ls = []

iteration

t=0

while(t<self.max_iteration):

while(t<self.max_iteration):

t+=1

pick a random point and add bias term

rand_index = np.random.randint(0,len(train_y))

x = np.append(train_x[rand_index],1)

y = train_y[rand_index]

x = np.array([x]).T

y = np.array([y]).T

forward pass

z = np.tanh((self.V.T).dot(x))

z = np.vstack((z,[1]))

h = expit((self.W.T).dot(z))

backward pass

```

if self.cost_function == "cross-entropy":
    delta_V,delta_W = self.Cross_entropy_delta_W(z,h,y,x,self.W)

    ### gradient descent update ###
    self.V = (1-self.regularizor)*self.V - self.learning_rate*delta_V
    self.W = (1-self.regularizor)*self.W - self.learning_rate*delta_W

    ### total training error and classification accuracy vs. iteration ###
    if t%1000==0:
        self.iteration_nums.append(t)
        if self.cost_function=="cross-entropy":
            error = self.cross_entropy_error(train_x,train_y,self.V,self.W)
            predict_labels = self.predict(train_x)
            true_labels = train_y.argmax(axis=1)
            accuracy = sum(predict_labels==true_labels)/len(true_labels)
            print(t, " : ",error,accuracy)
            self.error_ls.append(error)
            self.accuracy_ls.append(accuracy)
        if t%(train_x.shape[0])==0:
            k = t/(train_x.shape[0])
            self.learning_rate = (1.0*k/(k+1))*self.learning_rate

def Cross_entropy_delta_W(self,z,h,y,x,W):
    delta_W = z.dot((h-y).T)
    delta_V = x.dot(((W[0:-1].dot(h-y))*(1-(z[0:-1])**2)).T)
    return delta_V,delta_W

def cross_entropy_error(self,train_x,train_y,V,W):
    ### calculate total cross entropy error ###
    X = np.column_stack((train_x, np.ones(train_x.shape[0])))
    Z = np.tanh(X.dot(V))
    Z = np.column_stack((Z, np.ones(Z.shape[0])))
    H = 1/(np.exp(-Z.dot(W))+1)
    Cross_entropy_error = np.sum(-(train_y*np.log(H)+(1-train_y)*np.log(1-H)))
    return Cross_entropy_error

def predict(self,test_x):
    X = np.column_stack((test_x, np.ones(test_x.shape[0])))
    Z = np.tanh(X.dot(self.V))
    Z = np.column_stack((Z, np.ones(Z.shape[0])))
    H = Z.dot(self.W)
    predicted_labels = H.argmax(axis=1)
    return predicted_labels

```

```
#predict the test dataset#
final_NN = NN2(learning_rate=0.001,max_iteration=600000,cost_function="cross-
entropy",regularizer=0,n_hid=500)
final_NN.train(train_x,train_y)
test_y = final_NN.predict(test_x)
np.savetxt('letters_predict.csv', test_y, delimiter = ',')
```